
NN Synthesis Project

吳育丞 吳冠緯

Outline

- 1. Model Training
 - Data Augmentation
 - ReLU Upper Bound
 - Hyper Parameter (lr, batch size, optimizer ...)
 - Knowledge Distillation
- 2. Circuit Generation
 - Generate Matrix
 - Find Sharing
 - Construct Graph & MaxWeightMatching
- 3. FPGA Implementation (Bonus)
 - Simulation
 - Synthesis
 - Implementation

1. Model Training

Data Augmentation

```
transforms.ToTensor(),  
transforms.RandomRotation(5),  
transforms.RandomResizedCrop(28),  
transforms.Normalized((0.5,), (0.5,))
```

ReLU Upper Bound

`torch.nn.functional.relu6()` → limit upper bound to 6

[$\text{ReLU6}(x) = \min(\max(0, x), 6)$]

`torch.nn.functional.relu() + min(x, 7)` → limit upper bound to 7

4 bits are used on both upper bounds

Hyper Parameter

Default:

bit = 4

lr = 0.0005 (0.005, 0.00005 ...)

dropout = 0.35 (0.5 ...)

batch size = 256 (128 ...)

epoch = 100 (150 ...)

optimizer = Adam (RAdam, SGD ...)

Large Model Example

c1 (1, 3, k=3, s=2, p=0)

f1 (3*13*13, 64)

f2 (64, 32)

f3 (32, 10)

Testing accuracy of the quantized model: 0.9701

```
--- Number of parameter ---  
34958  
epoch:5, train acc:0.9141, valid acc:0.9061  
epoch:10, train acc:0.94134, valid acc:0.9355  
epoch:15, train acc:0.95994, valid acc:0.9555  
epoch:20, train acc:0.96434, valid acc:0.9565  
epoch:25, train acc:0.96898, valid acc:0.9603  
epoch:30, train acc:0.97014, valid acc:0.964  
epoch:35, train acc:0.9699, valid acc:0.9617  
epoch:40, train acc:0.97166, valid acc:0.9626  
epoch:45, train acc:0.97664, valid acc:0.9679  
epoch:50, train acc:0.97474, valid acc:0.9666  
epoch:55, train acc:0.97512, valid acc:0.9682  
epoch:60, train acc:0.97536, valid acc:0.9666  
epoch:65, train acc:0.97206, valid acc:0.9648  
epoch:70, train acc:0.9756, valid acc:0.9683  
epoch:75, train acc:0.97658, valid acc:0.9676  
epoch:80, train acc:0.97002, valid acc:0.9627  
epoch:85, train acc:0.97458, valid acc:0.965  
epoch:90, train acc:0.97388, valid acc:0.966  
epoch:95, train acc:0.97612, valid acc:0.9686  
epoch:100, train acc:0.9746, valid acc:0.9652  
best valid acc: 0.9686
```

Small Model Example

c1 (1, 3, k=11, s=3, p=0)

f1 (3*6*6, 10)

Testing accuracy of the quantized model: 0.953

```
--- Number of parameter ---  
1462  
epoch:5, train acc:0.90762, valid acc:0.9062  
epoch:10, train acc:0.92722, valid acc:0.9291  
epoch:15, train acc:0.9338, valid acc:0.9311  
epoch:20, train acc:0.94408, valid acc:0.9421  
epoch:25, train acc:0.94722, valid acc:0.9465  
epoch:30, train acc:0.94766, valid acc:0.9452  
epoch:35, train acc:0.94636, valid acc:0.9446  
epoch:40, train acc:0.95076, valid acc:0.9447  
epoch:45, train acc:0.9503, valid acc:0.95  
epoch:50, train acc:0.93756, valid acc:0.9405  
epoch:55, train acc:0.94084, valid acc:0.9386  
epoch:60, train acc:0.94088, valid acc:0.9404  
epoch:65, train acc:0.93826, valid acc:0.9335  
epoch:70, train acc:0.93, valid acc:0.9316  
epoch:75, train acc:0.94516, valid acc:0.9457  
epoch:80, train acc:0.94968, valid acc:0.9442  
epoch:85, train acc:0.94236, valid acc:0.9394  
epoch:90, train acc:0.92982, valid acc:0.9312  
epoch:95, train acc:0.94402, valid acc:0.9409  
epoch:100, train acc:0.9434, valid acc:0.9421  
best valid acc: 0.95
```


Different Modification: Model Architecture

fully connected NN model (4 layers):

256 / 128 / 64 / 10 : 98%

128 / 64 / 32 / 10 : 97%

64 / 48 / 16 / 10 : 97%

48 / 32 / 16 / 10 : 95%

Note that for 2 layers 256 / 10 can obtain 97% acc as well.

Different Modification: lr

change the learning rate of Adam optimizer:

batch size = 256

bit = 4

	lr = 0.005	lr = 0.0005	lr = 0.00005
best valid	0.9356	0.9566	0.9529
test accuracy	0.9347	0.9547	0.9538

Different Modification: bit

change the learning rate of Adam optimizer:

batch size = 256

lr = 0.0005

	bit = 4	bit = 6	bit = 8
best valid	0.9566	0.9776	0.9777
test accuracy	0.9547	0.9764	0.9764

Different Modification: batch size

change the learning rate of Adam optimizer:

bit = 4

lr = 0.0005

	size = 256	size = 128	size = 64
best valid	0.9566	0.9488	0.953
test accuracy	0.9574	0.9512	0.9593

Different Modification: ReLU upper bound

change the learning rate of Adam optimizer:

bit = 4

lr = 0.0005

batch size = 256

	limit = 5	limit = 6 (relu6)	limit = 7
best valid	0.9575	0.9628 (0.9566)	0.9536
test accuracy	0.9618	0.9634 (0.9574)	0.9535

Different Modification: bit quantization

Quantizing activations (a) and weights (w) to 4-bit and 3-bit

default (case I): conv (1, 3, 3, 2) / fc (3*13*13, 32, 32, 10)

case II: fc (3*13*13, 80, 32, 10) / case III: conv (1, 5, 3, 2)

	a4 w4 (I)	a4 w3 (I)	a3 w3 (I)	a3 w3 (II)	a3 w3 (III)
best valid	95.99%	96.22%	93.86%	94.92%	95.68%
test accu.	96.21%	96.21%	93.61%	95.42%	95.32%
AIG accu.	95.73%	86.93%	84.8%	95.1%	93.77%
AND gates #	516492	309698	245716	323941	325484

Different Modification: Knowledge Distillation

lr = 0.002

batch size = 256

teacher model: conv (1, 8, 3, 2, 0) / fc (512, 128, 10)


use learning rate schedule: ExponentialLR with gamma = 0.95

without KD (with KD)	4-bit teacher	4-bit student	3-bit teacher	3-bit student ($\alpha=0.9$)	3-bit student ($\alpha=0.9-0.1$)
best valid	98.56	96.24(95.66)	98	95.15(93.86)	94.08(93.86)
test accu.	98.36	96.12(95.47)	97.91	95.06(93.61)	94.49(93.61)

2. Circuit Generation

Y.-S. Huang, J.-H. R. Jiang, and A. Mishchenko,
"Quantized neural network synthesis for direct logic circuit implementation", Proc. IWLS'21.

Generate Matrix



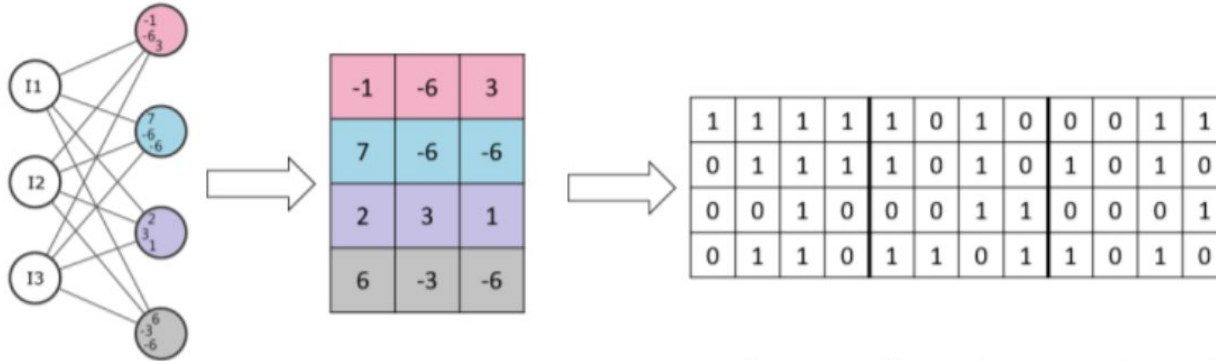
$$O_1 = -I_1 - 6I_2 + 3I_3$$

$$O_2 = 7I_1 - 6I_2 - 6I_3$$

$$O_3 = 2I_1 + 3I_2 + I_3$$


$$O_4 = 6I_1 - 3I_2 - 6I_3$$

Generate Matrix



$$\begin{aligned}
 O_1 &= -(I_1 \ll 3) + (I_1 \ll 2) + (I_1 \ll 1) + I_1 - \\
 &\quad (I_2 \ll 3) + (I_2 \ll 1) + (I_3 \ll 1) + I_3 \\
 O_2 &= (I_1 \ll 2) + (I_1 \ll 1) + I_1 - (I_2 \ll 3) + (I_2 \ll 1) - \\
 &\quad (I_3 \ll 3) + (I_3 \ll 1) \\
 O_3 &= (I_1 \ll 1) + (I_2 \ll 1) + I_2 + I_3 \\
 O_4 &= (I_1 \ll 2) + (I_1 \ll 1) - (I_2 \ll 3) + (I_2 \ll 2) + I_2 - \\
 &\quad (I_3 \ll 3) + (I_3 \ll 1)
 \end{aligned}$$

Booth Matrix



$$O_1 = -I_1 - (I_2 \ll 3) + (I_2 \ll 1) + (I_3 \ll 2) - I_3$$



$$O_2 = (I_1 \ll 3) - I_1 - (I_2 \ll 3) + (I_2 \ll 1) - (I_3 \ll 3) + (I_3 \ll 1)$$

$$O_3 = (I_1 \ll 1) + (I_2 \ll 2) - I_2 + I_3$$

$$O_4 = (I_1 \ll 3) - (I_1 \ll 1) - (I_2 \ll 2) + I_2 - (I_3 \ll 3) + (I_3 \ll 1)$$

Find Sharing

0	0	0	-1	-1	0	1	0	0	1	0	-1
1	0	0	-1	-1	0	1	0	-1	0	1	0
0	0	1	0	0	1	0	-1	0	0	0	1
1	0	-1	0	0	-1	0	1	-1	0	1	0



$$O_1 = -I_1 - (I_2 \ll 3) + (I_2 \ll 1) + (I_3 \ll 2) - I_3$$



$$O_2 = (I_1 \ll 3) - I_1 - (I_2 \ll 3) + (I_2 \ll 1) - (I_3 \ll 3) + (I_3 \ll 1)$$

$$O_3 = (I_1 \ll 1) + (I_2 \ll 2) - I_2 + I_3$$

$$O_4 = (I_1 \ll 3) - (I_1 \ll 1) - (I_2 \ll 2) + I_2 - (I_3 \ll 3) + (I_3 \ll 1)$$

Find Sharing

0	0	0	-1	-1	0	1	0	0	1	0	-1
1	0	0	-1	-1	0	1	0	-1	0	1	0
0	0	1	0	0	1	0	-1	0	0	0	1
1	0	-1	0	0	-1	0	1	-1	0	1	0



$$O_1 = -I_1 - (I_2 \ll 3) + (I_2 \ll 1) + (I_3 \ll 2) - I_3$$

$$O_2 = (I_1 \ll 3) - I_1 - (I_2 \ll 3) + (I_2 \ll 1) - (I_3 \ll 3) + (I_3 \ll 1)$$

$$O_3 = (I_1 \ll 1) + (I_2 \ll 2) - I_2 + I_3$$


$$O_4 = (I_1 \ll 3) - (I_1 \ll 1) - (I_2 \ll 2) + I_2 - (I_3 \ll 3) + (I_3 \ll 1)$$

Find Sharing

```
Algorithm FindMaxPairing ( $W_B$ ,  $unSharedCols$ )  
   $G(V, E) = \text{ConstructGraph}(W_B, unSharedCols)$   
   $pairs = \text{MaxWeightMatching}(V, E)$   
   $result = []$   
  foreach  $pair$  in  $pairs$ :  
     $s^{same}, s^{diff} = \text{PairRows}(pair)$   
    if gain ( $s^{same}$ ) > 0:  
      append  $s^{same}$  to  $result$   
    if gain ( $s^{diff}$ ) > 0:  
      append  $s^{diff}$  to  $result$   
  return  $result$ 
```

Construct Graph

0	0	0	-1	-1	0	1	0	0	1	0	-1
1	0	0	-1	-1	0	1	0	-1	0	1	0
0	0	1	0	0	1	0	-1	0	0	0	1
1	0	-1	0	0	-1	0	1	-1	0	1	0



MaxWeightMatching

Z. Galil, "Efficient algorithms for finding maximum matching in graphs," *ACM Comput. Surv.*, vol. 18, no. 1, pp. 23–38, 1986.



Overall Algorithm

Algorithm 1 FindSharing

Input: An $m \times nb$ Booth matrix W_B

Output: A list of shared terms


```
1:  $C := \{0, \dots, nb - 1\}$ ;  
2:  $unSharedCols :=$  an empty hash table;  
3: for  $r = 0$  to  $m - 1$  do  
4:    $unSharedCols[r] = C$ ;  
5:  $result := []$ ;  
6:  $sharings := FindMaxPairing(W_B, unSharedCols)$ ;  
7: while  $size(sharings) > 0$  do  
8:   append elements in  $sharings$  to  $result$ ;  
9:    $unSharedCols := refine(unSharedCols, sharings)$ ;  
10:   $sharings := FindMaxPairing(W_B, unSharedCols)$ ;  
11: return  $result$ ;
```

3. FPGA Implementation (Bonus)

Example Result

c1 (1, 2, k=3, s=2, p=0)

f1 (2*13*13, 10)



New Project Summary

- i** A new RTL project named 'Presentation_Example' will be created.
- i** 5 source files will be added.
- i** 1 constraints file will be added.
- i** The default part and product family for the new project:
 - Default Board: Nexys Video**
 - Default Part: xc7a200tsbg484-1
 - Product: Artix-7
 - Family: Artix-7
 - Package: sbg484
 - Speed Grade: -1

Simulation

```
pred:0000100000
correct
=====
correct count:      9285
=====
index:      9999 true label:  6

pred:0001000000
correct
=====
correct count:      9286
$finish called at time : 1400105 ns : File "C:/Users/User/Desktop/eda_project/W3/lab/files/sim/tb_unpacked.sv" Line 59
run: Time (s): cpu = 00:00:16 ; elapsed = 00:00:12 . Memory (MB): peak = 1395.102 ; gain = 36.078
```

Synthesis: Time

✓ synth_1 (active)	constrs_1	synth_design Complete!									69184	6716	0.0
▶ impl_1	constrs_1	Not started											
Out-of-Context Module Runs													
✓ blk_mem_gen_0_synth_1	blk_mem_gen_0	synth_design Complete!									0	0	174.5

Setup		Hold		Pulse Width	
Worst Negative Slack (WNS): 3.377 ns		Worst Hold Slack (WHS): 0.139 ns		Worst Pulse Width Slack (WPWS): 9.500	
Total Negative Slack (TNS):	0.000 ns	Total Hold Slack (THS):	0.000 ns	Total Pulse Width Negative Slack (TPWS):	0.000
Number of Failing Endpoints:	0	Number of Failing Endpoints:	0	Number of Failing Endpoints:	0
Total Number of Endpoints:	15368	Total Number of Endpoints:	15368	Total Number of Endpoints:	7067
All user specified timing constraints are met.					


Synthesis: Power

Power estimation from Synthesized netlist. Activity derived from constraints files, simulation files or vectorless analysis. Note: these early estimates can change after implementation.

Total On-Chip Power:	0.998 W
Design Power Budget:	Not Specified
Power Budget Margin:	N/A
Junction Temperature:	28.3°C
Thermal Margin:	56.7°C (16.7 W)
Effective θ_{JA} :	3.3°C/W
Power supplied to off-chip devices:	0 W
Confidence level:	Medium

[Launch Power Constraint Advisor](#) to find and fix invalid switching activity

On-Chip Power



Implementation: Time

Name	Constraints	Status	WNS	TNS	WHS	THS	TPWS	Total Power	Failed Routes	LUT	FF
✓ synth_1 (active)	constrs_1	synth_design Complete!								69184	6716
✓ impl_1	constrs_1	route_design Complete!	0.471	0.000	0.066	0.000	0.000	0.659	0	57876	6717
Out-of-Context Module Runs											
✓ blk_mem_gen_0_synth_1	blk_mem_gen_0	synth_design Complete!								0	0

BRAM	URAM	DSP	Start	Elapsed	Run Strategy	Report Strategy
0.0	0	0	6/16/22, 11:27 PM	00:14:51	Flow_AlternateRoutability (Vivado Synthesis 2021)	Vivado Synthesis Default Reports (Vivado Synthesis 2021)
174.5	0	0	6/16/22, 11:44 PM	00:29:06	Congestion_SpreadLogic_medium (Vivado Implementation 2021)	Vivado Implementation Default Reports (Vivado Implementation 2021)
174.5	0	0	6/16/22, 11:19 PM	00:05:32	Vivado Synthesis Defaults (Vivado Synthesis 2021)	Vivado Synthesis Default Reports (Vivado Synthesis 2021)


Setup	Hold	Pulse Width
Worst Negative Slack (WNS): 0.471 ns	Worst Hold Slack (WHS): 0.066 ns	Worst Pulse Width Slack (WPWS): 9.500 ns
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): 0.000 ns
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: 0
Total Number of Endpoints: 15719	Total Number of Endpoints: 15719	Total Number of Endpoints: 7068
All user specified timing constraints are met.		

Implementation: Power

Power analysis from Implemented netlist. Activity derived from constraints files, simulation files or vectorless analysis.

Total On-Chip Power:	0.659 W
Design Power Budget:	Not Specified
Power Budget Margin:	N/A
Junction Temperature:	27.2°C
Thermal Margin:	57.8°C (17.1 W)
Effective θ_{JA}:	3.3°C/W
Power supplied to off-chip devices:	0 W
Confidence level:	Medium

[Launch Power Constraint Advisor](#) to find and fix invalid switching activity



Q&A